

Engineering AI-Powered Chatbots with RAG and LLMs

Engineering AI-Powered Chatbots with RAG and LLMs

👤 Done by: Malak Albahri

✉ Email: malakbaderalbahri@gmail.com

🔗 LinkedIn: [Malak Albahri](#)

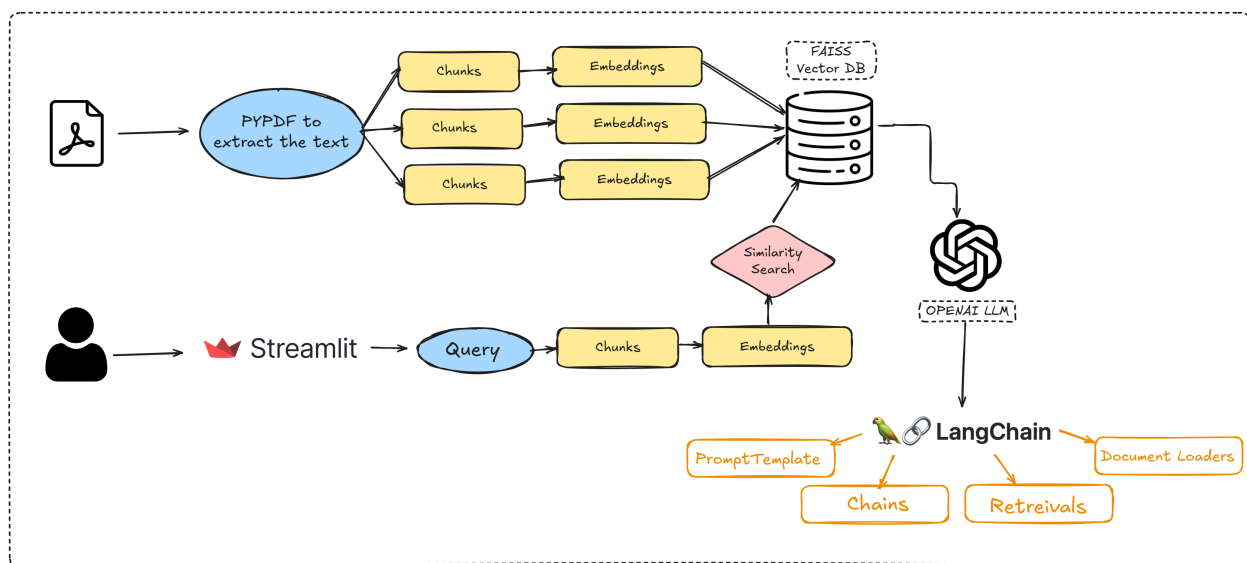
🐙 GitHub: [MALAKBADEROO](#)

Welcome to this workshop summary on Building an AI chatbot using **Retrieval-Augmented Generation (RAG)**! 🚀

- **What is RAG?** A powerful AI technique that combines **retrieval** (fetching relevant data) with **generation** (LLM-powered responses).
- **Why use it?** To enhance responses by grounding them in real data.
- **What's an LLM?** A **Large Language Model** like OpenAI's GPT that can generate human-like text.

Let's dive in! 🔍

RAG Structure



? How do we extract PDFs?

✓ We use **LangChain document loaders**, which work on top of **PyPDF** to read and process PDFs seamlessly.

? Why do we need chunking?

🔧 Context windows in LLMs are limited! Chunking helps break long texts into digestible pieces, making retrieval more efficient.

? What are embeddings?

🧠 Converting human language into vectors! Just like **linear algebra**, we map words into numerical space. We use **OpenAI embeddings** for this.

? What's a Vector Database? Why FAISS?

🔍 A database optimized for **fast similarity searches**! **FAISS** is easy to use and efficient.

? How do users interact with our system?

🖥️ **Streamlit UI**! A simple yet powerful way to create interactive web apps in Python.

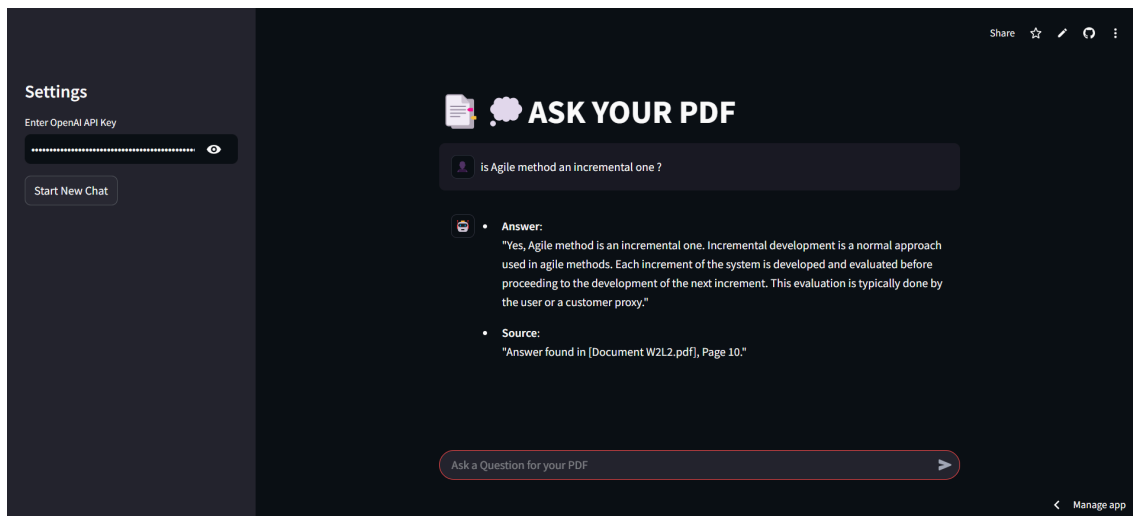
? What does LangChain offer us?

🔧 **Prompt templates, chains, retrievers, document loaders**—all built-in!

? How do we retrieve relevant information?

🔍 By default, **similarity search** is used to find the closest matches in our vector database.

Final Look - The Streamlit UI



🔧 1. Setting Up Python

📌 Step 1: Download & Install Python

🔗 [Download Python](#)

Ensure you check the option "Add Python to PATH" during installation.

📌 Step 2: Verify Python Installation

Run the following command:

OS	Command
Windows	<code>python --version</code>
macOS	<code>python3 --version</code>

🌐 2. Creating a Virtual Environment

📌 Step 1: Create a Virtual Environment

OS	Command
Windows	<code>python -m venv venv</code>
macOS	<code>python3 -m venv venv</code>

📌 Step 2: Activate the Virtual Environment

OS	Command
Windows	<code>venv\Scripts\activate</code>
macOS	<code>source venv/bin/activate</code>

📦 3. Installing Dependencies

📌 Step 1: Create a `requirements.txt` File

Inside your project folder, create a file called `requirements.txt` and add the following:

```
langchain
langchain_openai
langchain_community
streamlit
python_dotenv
ipykernel
pypdf
faiss-cpu
```

📌 Step 2: Install Dependencies

OS	Command
Windows	<code>pip install -r requirements.txt</code>
macOS	<code>pip3 install -r requirements.txt</code>

📁 4. Organizing Project Structure

Your project should be structured as follows:

```
/chatbot-project
| — /venv          # Virtual Environment
| — /data          # Folder to store PDFs
| — /faiss_index   # Folder to store embeddings (Auto-created)
| — requirements.txt # Dependencies
| — app.py         # Streamlit UI
| — rag.py         # RAG pipeline
```

5. Writing `app.py` and `rag.py`

The core of our project consists of two key files:

 `rag.py` – Handles the Retrieval-Augmented Generation (RAG) pipeline.

 `app.py` – A Streamlit-powered UI for user interaction.

 **Good News!** The complete code has been **uploaded to GitHub**, so you can check it out and use it directly:

 **GitHub Repository:** [githubRepo](#)



 You can download or clone the repo and use `rag.py` and `app.py` as a reference to build your own chatbot! 

6. Running the Chatbot UI

 **Run Streamlit**

OS	Command
Windows	<code>streamlit run app.py</code>
macOS	<code>streamlit run app.py</code>

7. OpenAi API Key Extraction

-  **Get Your API Key:** [OpenAI API Keys](#)
-  **Watch This Guide:** [YouTube Tutorial](#)

Pushing Code to GitHub - Step by Step

1 Install Git

- Download Git from [here](#).
- During installation, **check "Add Git to PATH"** before clicking Next.

2 Create a GitHub Repository

- Go to [GitHub](#) and create a new repo.
- Name it, add a description, set it to **Public**, and click **Create**.

3 Set Up Git in VS Code

- Copy the GitHub setup commands after repo creation.
- Open **VS Code Terminal** (**Ctrl + ~**) and paste the commands.
- Refresh GitHub, and you'll see the repo with only a README file.

4 Add a `.gitignore` File

- On GitHub, go to **Add file** → **Create new file**.
- Name it `.gitignore`, select **Python**, and commit changes.
- In VS Code, run:

```
git pull
```

- `.gitignore` prevents pushing unnecessary/sensitive files.

5 Push Your Code to GitHub

Run these commands in VS Code Terminal:

```
git add .  
git commit -m "Initial commit"  
git push -u origin main
```

- Refresh GitHub → Your code is now uploaded. 🎉

🌐 Deploying the Streamlit App

- Run the App in locally using `streamlit run app.py`
- Click the **Deploy** button and set your desired web URL name.



Explore My App: It features **Software Engineering course content** – check it out here: comp3401chatbot.streamlit.app

📚 Resources & Documentation

- [OpenAI API Docs](#)
- [LangChain Docs](#)
- [Streamlit Docs](#)